

```
In [1]: import nltk, random
#function returning the last letter of a word
def gender_features(word):
    return {'last_letter': word[-1]}
```

```
In [86]: gender_features('Shrek')
```

```
Out[86]: {'last_letter': 'k'}
```

```
In [87]: from nltk.corpus import names# set of male and female names:how to do our own sets?
labeled_names = [(name, 'male') for name in names.words('male.txt')] + [(name, 'female') for name in names.words('female.txt')]
```

```
In [88]: from sklearn.model_selection import train_test_split
featuresets = [(gender_features(n), gender) for (n, gender) in labeled_names]
train_set, test_set = train_test_split(featuresets, shuffle=True, test_size=0.2)
print(len(train_set))
print(len(test_set))
```

```
6355
1589
```

```
In [89]: classifier = nltk.NaiveBayesClassifier.train(train_set)# we need other steps like feature selection etc...
print(nltk.classify.accuracy(classifier, test_set))
```

```
0.7595972309628697
```

```
In [90]: print(classifier.classify(gender_features('Alan')))
print(classifier.classify(gender_features('Stacey')))
```

```
male
female
```

```
In [91]: classifier.show_most_informative_features(5)
```

```
Most Informative Features
  last_letter = 'k'           male : female =   38.0 : 1.0
  last_letter = 'a'           female : male   =   35.8 : 1.0
  last_letter = 'f'           male : female =   19.6 : 1.0
  last_letter = 'v'           male : female =   16.2 : 1.0
  last_letter = 'd'           male : female =    9.9 : 1.0
```

```
In [11]: from nltk.classify import apply_features
train_set, test_set = train_test_split(labeled_names, shuffle=True, test_size=.2)
train_set = apply_features(gender_features, train_set)
test_set = apply_features(gender_features, test_set)
```

## 1.2. Choosing the Right Features

```
In [94]: def gender_features2(name):
    features = {}
    features["first_letter"] = (name[0].lower() in 'aeiou')
    features["last_letter"] = (name[-1].lower() in 'aeiou')
    for letter in 'abcdefghijklmnopqrstuvwxyz':
        features["count({})".format(letter)] = name.lower().count(letter)
        features["has({})".format(letter)] = (letter in name.lower())
    return features
```

```
In [95]: gender_features2('Walter')
```

```
Out[95]: {'count(a)': 1,
'count(b)': 0,
'count(c)': 0,
'count(d)': 0,
'count(e)': 1,
'count(f)': 0,
'count(g)': 0,
'count(h)': 0,
'count(i)': 0,
'count(j)': 0,
'count(k)': 0,
'count(l)': 1,
'count(m)': 0,
'count(n)': 0,
'count(o)': 0,
'count(p)': 0,
'count(q)': 0,
'count(r)': 1,
'count(s)': 0,
'count(t)': 1,
'count(u)': 0,
'count(v)': 0,
'count(w)': 1,
'count(x)': 0,
'count(y)': 0,
'count(z)': 0,
'first_letter': 'w',
'has(a)': True,
'has(b)': False,
'has(c)': False,
'has(d)': False,
'has(e)': True,
'has(f)': False,
'has(g)': False,
'has(h)': False,
'has(i)': False,
'has(j)': False,
'has(k)': False,
'has(l)': True,
'has(m)': False,
'has(n)': False,
'has(o)': False,
'has(p)': False,
'has(q)': False,
'has(r)': True,
'has(s)': False,
'has(t)': True,
'has(u)': False,
'has(v)': False,
'has(w)': True,
'has(x)': False,
'has(y)': False,
'has(z)': False,
'last_letter': 'r'}
```

```
In [96]: featuresets = [(gender_features2(n), gender) for (n, gender) in labeled_names]
train_set, test_set = train_test_split(featuresets, shuffle=True, test_size=.2)
classifier = nltk.NaiveBayesClassifier.train(train_set)
print(nltk.classify.accuracy(classifier, test_set))
```

```
0.7627438640654499
```

```
In [97]: train_names, test_names = train_test_split(labeled_names, shuffle=True, test_size=.3)
devtest_names, test_names = train_test_split(test_names, shuffle=True, test_size=.3333)
```

```
In [98]: [print(len(s)) for s in [train_names, devtest_names, test_names]]
```

```
5560
1589
795
```

```
Out[98]: [None, None, None]
```

```
In [99]: train_set = [(gender_features(n), gender) for (n, gender) in train_names]
devtest_set = [(gender_features(n), gender) for (n, gender) in devtest_names]
test_set = [(gender_features(n), gender) for (n, gender) in test_names]
classifier = nltk.NaiveBayesClassifier.train(train_set)
print(nltk.classify.accuracy(classifier, devtest_set))
```

```
0.7602265575833858
```

```
In [100]: errors = []
for (name, tag) in devtest_names:
    guess = classifier.classify(gender_features(name))
    if guess != tag:
        errors.append( (tag, guess, name) )
```

```
In [101]: for (tag, guess, name) in sorted(errors):
    print('correct={:8} guess={:8s} name={:30}'.format(tag, guess, name))
```

```
correct=female guess=male name=Adriaens
correct=female guess=male name=Aigneis
correct=female guess=male name=Alex
correct=female guess=male name=Alexis
correct=female guess=male name=Alis
correct=female guess=male name=Alisun
correct=female guess=male name=Allien
correct=female guess=male name=Allyson
correct=female guess=male name=Alyson
correct=female guess=male name=Anne-Mar
correct=female guess=male name=Ashlien
correct=female guess=male name=Betteann
correct=female guess=male name=Bev
correct=female guess=male name=Blair
correct=female guess=male name=Bo
correct=female guess=male name=Brear
correct=female guess=male name=Britt
correct=female guess=male name=Brittan
correct=female guess=male name=Brooks
```

```
In [103]: def gender_features(word):
          return {'suffix1': word
                  'suffix2': word
```

```
In [104]: train_set = [(gender_features(n), gender) for (n, gender) in train_names]
devtest_set = [(gender_features(n), gender) for (n, gender) in devtest_names]
classifier = nltk.NaiveBayesClassifier.train(train_set)
print(nltk.classify.accuracy(classifier, devtest_set))
#we need more on feature selection, chisquare etc...accuracy chart

0.7853996224040277
```

### 1.3 Document Classification

```
In [108]: from nltk.corpus import movie_reviews
documents = [(list(movie_reviews.words(fileid)), category)
              for category in movie_reviews.categories()
              for fileid in movie_reviews.fileids(category)]
random.shuffle(documents)
```

```
In [110]: all_words = nltk.FreqDist(w.lower() for w in movie_reviews.words())
word_features = list(all_words)[:2000]
def document_features(document):
    document_words = set(document)
    features = {}
    for word in word_features:
        features['contains({})'.format(word)] = (word in document_words)
    return features
```

```
In [111]: featuresets = [(document_features(d), c) for (d,c) in documents]
train_set, test_set = train_test_split(featuresets, shuffle=True, test_size=.2)
# train_set, test_set = featuresets[100:], featuresets[:100]
classifier = nltk.NaiveBayesClassifier.train(train_set)
```

```
In [116]: print(nltk.classify.accuracy(classifier, test_set))
classifier.show_most_informative_features(10)

0.6725
Most Informative Features
contains(hudson) = True          neg : pos = 7.8 : 1.0
contains(et) = True             pos : neg = 7.6 : 1.0
contains(taxi) = True           pos : neg = 7.6 : 1.0
contains(stupidity) = True      neg : pos = 7.4 : 1.0
contains(debate) = True        pos : neg = 7.3 : 1.0
contains(ivy) = True            neg : pos = 6.4 : 1.0
contains(confidential) = True   pos : neg = 6.3 : 1.0
contains(whore) = True          neg : pos = 5.7 : 1.0
contains(lang) = True           pos : neg = 5.6 : 1.0
contains(freely) = True         pos : neg = 5.6 : 1.0
```

### 1.4 Part-of-Speech Tagging

```
In [3]: from nltk.corpus import brown
suffix_fdlist = nltk.FreqDist()
for word in brown.words():
    word = word.lower()
    suffix_fdlist[word[-1:]] += 1
    suffix_fdlist[word[-2:]] += 1
    suffix_fdlist[word[-3:]] += 1
```

```
In [4]: common_suffixes = [suffix for (suffix, count) in suffix_fdlist.most_common(100)]
print(common_suffixes)
```

```
['e', ',', '.', 's', 'd', 't', 'he', 'n', 'a', 'of', 'the', 'y', 'r', 'to', 'in', 'f', 'o', 'ed', 'nd', 'is', 'on', 'l', 'g', 'and', 'ng', 'en', 'as', 'ing', 'h', 'at', 'es', 'on', 'r',
'e', 'it', 'an', 'm', 'i', 'ly', 'ion', 'en', 'al', '?', 'nt', 'be', 'hat', 'st', 'his', 'th', 'll', 'le', 'ce', 'by', 'ts', 'me', 've', 'se', 'ut', 'was', 'for',
'ent', 'ch', 'k', 'w', 'ld', 'rs', 'ted', 'ere', 'her', 'ne', 'ns', 'ith', 'ad', 'ry', ')', '(', 'te', '--', 'ay', 'ty', 'ot', 'p', 'nce', 's', 'ter', 'om', 'ss', ':', 'we', 'ar',
'e', 'c', 'ers', 'uld', 'had', 'so', 'ey']
```

```
In [9]: def pos_features(word):
          features = {}
          for suffix in common_suffixes:
              features['endswith({})'.format(suffix)] = word.lower().endswith(suffix)
          return features
```

```
In [ ]: tagged_words = brown.tagged_words(categories='news')
featuresets = [(pos_features(n), g) for (n,g) in tagged_words]
```

```
In [7]: # size = int(len(featuresets) * 0.1)
# train_set, test_set = featuresets[size:], featuresets[:size]
train_set, test_set = train_test_split(featuresets, shuffle=False, test_size=.1)

-----
NameError                                Traceback (most recent call last)
<ipython-input-7-12cc5a8dda35> in <module>()
      1 # size = int(len(featuresets) * 0.1)
      2 # train_set, test_set = featuresets[size:], featuresets[:size]
----> 3 train_set, test_set = train_test_split(featuresets, shuffle=False, test_size=.1)

NameError: name 'train_test_split' is not defined
```

```
In [124]: classifier = nltk.DecisionTreeClassifier.train(train_set)
nltk.classify.accuracy(classifier, test_set)

-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-124-0c836654950f> in <module>()
----> 1 classifier = nltk.DecisionTreeClassifier.train(train_set)
      2 nltk.classify.accuracy(classifier, test_set)

/usr/local/lib/python3.5/dist-packages/nltk/classify/decisiontree.py in train(labeled_featuresets, entropy_cutoff, depth_cutoff, support_cutoff, binary, feature_values, verbose)
    159     # Refine the stump.
    160     tree.refine(labeled_featuresets, entropy_cutoff, depth_cutoff-1,
-> 161               support_cutoff, binary, feature_values, verbose)
    162
    163     # Return it

/usr/local/lib/python3.5/dist-packages/nltk/classify/decisiontree.py in refine(self, labeled_featuresets, entropy_cutoff, depth_cutoff, support_cutoff, binary, feature_values, verbose)
    201         self._decisions[fval] = DecisionTreeClassifier.train(
    202             fval_featuresets, entropy_cutoff, depth_cutoff,
-> 203             support_cutoff, binary, feature_values, verbose)
    204         if self._default is not None:
    205             # If there is a default class, then we can return it.

In [ ]: classifier.classify(pos_features('dog'))

In [34]: print(classifier.pseudocode(depth=4))
```

## 1.5 Exploiting Context

```
In [36]: def pos_features(sentence, i):
          features = {"suffix(1)": sentence[i][-1:],
                    "suffix(2)": sentence[i][-2:],
                    "suffix(3)": sentence[i][-3:]}

          if i == 0:
              features["prev-word"] = "<START>"
          else:
              features["prev-word"] = sentence[i-1]
          return features

In [37]: pos_features(brown.sents()[0], 8)

Out[37]: {'prev-word': 'an', 'suffix(1)': 'n', 'suffix(2)': 'on', 'suffix(3)': 'ion'}
```

```
In [38]: tagged_sents = brown.tagged_sents(categories='news')
featuresets = []
for tagged_sent in tagged_sents:
    untagged_sent = nltk.tag.untag(tagged_sent)
    for i, (word, tag) in enumerate(untagged_sent):
        featuresets.append( (pos_features(untagged_sent, i), tag) )

In [39]: # size = int(len(featuresets) * 0.1)
# train_set, test_set = featuresets[:size], featuresets[size:]
train_set, test_set = train_test_split(featuresets, shuffle=False, test_size=.1)
classifier = nltk.NaiveBayesClassifier.train(train_set)
nltk.classify.accuracy(classifier, test_set)

Out[39]: 0.7771479713603818
```

```
In [40]: def pos_features(sentence, i, history):
          features = {"suffix(1)": sentence[i][-1:],
                    "suffix(2)": sentence[i][-2:],
                    "suffix(3)": sentence[i][-3:]}

          if i == 0:
              features["prev-word"] = "<START>"
              features["prev-tag"] = "<START>"
          else:
              features["prev-word"] = sentence[i-1]
              features["prev-tag"] = history[i-1]
          return features

class ConsecutivePosTagger(nltk.TaggerI):
    def __init__(self, train_sents):
        train_set = []
        for tagged_sent in train_sents:
            untagged_sent = nltk.tag.untag(tagged_sent)
            history = []
            for i, (word, tag) in enumerate(untagged_sent):
                featureset = pos_features(untagged_sent, i, history)
                train_set.append( (featureset, tag) )
            history.append(tag)
        self.classifier = nltk.NaiveBayesClassifier.train(train_set)

    def tag(self, sentence):
        history = []
        for i, word in enumerate(sentence):
            featureset = pos_features(sentence, i, history)
            tag = self.classifier.classify(featureset)
            history.append(tag)
        return zip(sentence, history)
```

```
In [41]: tagged_sents = brown.tagged_sents(categories='news')
# size = int(len(tagged_sents) * 0.1)
# train_sents, test_sents = tagged_sents[:size], tagged_sents[size:]
train_sents, test_sents = train_test_split(tagged_sents, shuffle=False, test_size=.1)
tagger = ConsecutivePosTagger(train_sents)
print(tagger.evaluate(test_sents))
0.79796012981
```

0.7870028904614771

Out[41]: 0.79796012981

In [ ]: